

NGUYEN VO AN VI

Ho Chi Minh City, Vietnam | +84 362 699 765 | anvi23052006@gmail.com
[LinkedIn](#) | [GitHub](#) | [Kaggle](#) | [Portfolio](#)

PROFILE

Computer Science undergraduate at UIT focused on applied artificial intelligence and machine learning. Curious and self-driven, with hands-on experience in time-series forecasting, data-centric model development, information retrieval, and deployed RAG applications; seeking an AI/ML Engineering or Data Science internship in Ho Chi Minh City.

EDUCATION

University of Information Technology (UIT), VNU-HCM

B.Sc. in Computer Science

GPA: 8.61/10 | English Proficiency: VNU-EPT B2.1

Ho Chi Minh City, Vietnam

Sep 2024 - Expected 2028

TECHNICAL SKILLS

Languages: Python, C++, SQL, JavaScript

Machine Learning: PyTorch, scikit-learn, LightGBM, XGBoost, pandas, Polars

RAG & Retrieval: Gemini API, FAISS, BM25, embeddings, PEFT/LoRA

Development: FastAPI, React, Vite, REST APIs, Redis/Upstash, Git, Vercel, Render

SELECTED PROJECTS

UIT Knowledge Retrieval Chatbot | Gemini, FAISS, FastAPI, React

May - Jun 2026

AI Engineer | 3-member team | [GitHub](#) | [Live Demo](#)

- Built the Vietnamese RAG pipeline using Gemini embeddings and FAISS over 274 curated chunks from 84 official UIT webpages, with separate university-wide and department-specific knowledge indexes.
- Tuned L2 similarity-threshold filtering to reject low-relevance retrievals and improve the quality of context passed to Gemini.
- Implemented Gemini API rate-limit handling and Server-Sent Events streaming to maintain stable requests and progressively display generated responses.
- Developed the JavaScript/React interface, integrated it with FastAPI, and deployed the frontend on Vercel and backend on Render; won Third Prize among 15 DEV-O Hackathon finalist teams.

Hedge Fund Time-Series Forecasting | Python, LightGBM, Polars

Apr 2026

Technical Lead, Machine Learning | 5-member team | Ranked 7th/982 teams | [Leaderboard](#)

- Led the final LightGBM forecasting pipeline that ranked 7th out of 982 teams on the private leaderboard.
- Designed a 3-fold walk-forward validation scheme to preserve temporal order, prevent future-data leakage, and guide reliable offline model selection.
- Applied median imputation and percentile clipping with a 10% buffer, then benchmarked LightGBM, XGBoost, CatBoost, and an ensemble; retained LightGBM after the additional models degraded performance.
- Replaced an OOM-prone pandas workflow with Polars lazy execution, data-type optimization, and 200K-row batch inference, enabling stable processing of approximately 2 GB of data.

SELECTED COMPETITION EXPERIENCE

CLEF 2026 CheckThat! Lab - Scientific Document Retrieval

May 2026

Retrieval Research Contributor | 4-member team | Ranked 12th/42 teams

- Fine-tuned e5-large-v2 and BGE dense retrievers with PEFT/LoRA and contrastive learning; evaluated BM25-dense hybrid search, HyDE, hard-negative mining, and listwise reranking.
- Designed document-chunking experiments and contributed the Methodology, Dataset, and Experimental Setup sections to the team scientific report.

EXACT 2026 International Data Science Competition

Jun 2026

Neuro-Symbolic AI Contributor | 4-member team | Top 28/180 teams, Round 2

- Developed a Qwen-7B/8B and Z3 neuro-symbolic pipeline that translated natural-language reasoning problems into SMT constraints and verified candidate answers through satisfiability checks.
- Implemented robust JSON and logic parsing, syntax-error recovery, iterative hypothesis testing, and fallback handling for LLM-generated Z3 code across multiple-choice and yes/no tasks.

ADDITIONAL AWARD

Top 5 Finalist and Consolation Prize, Mastering IT 2026 - HCMUTE (56 qualifying teams)

May 2026